

# Chapter 1. Introducing C# and .NET Core

---

C# is a general-purpose, type-safe, object-oriented programming language. The goal of the language is programmer productivity. To this end, C# balances simplicity, expressiveness, and performance. The chief architect of the language since its first version is Anders Hejlsberg (creator of Turbo Pascal and architect of Delphi). The C# language is platform neutral and works with a range of platform-specific frameworks.

## Object Orientation

C# is a rich implementation of the object-orientation paradigm, which includes *encapsulation*, *inheritance*, and *polymorphism*.

Encapsulation means creating a boundary around an *object* to separate its external (public) behavior from its internal (private) implementation details. Following are the distinctive features of C# from an object-oriented perspective:

### Unified type system

The fundamental building block in C# is an encapsulated unit of data and functions called a *type*. C# has a *unified type system* in which all types ultimately share a common base type. This means that all types, whether they represent business objects or are primitive types such as numbers, share the same basic

functionality. For example, an instance of any type can be converted to a string by calling its `ToString` method.

## Classes and interfaces

In a traditional object-oriented paradigm, the only kind of type is a class. In C#, there are several other kinds of types, one of which is an *interface*. An interface is like a class that cannot hold data. This means that it can define only *behavior* (and not *state*), which allows for multiple inheritance as well as a separation between specification and implementation.

## Properties, methods, and events

In the pure object-oriented paradigm, all functions are *methods*. In C#, methods are only one kind of *function member*, which also includes *properties* and *events* (there are others, too). Properties are function members that encapsulate a piece of an object's state, such as a button's color or a label's text. Events are function members that simplify acting on object state changes.

Although C# is primarily an object-oriented language, it also borrows from the *functional programming* paradigm; specifically:

## Functions can be treated as values

Using *delegates*, C# allows functions to be passed as values to and from other functions.

## C# supports patterns for purity

Core to functional programming is avoiding the use of variables whose values change, in favor of declarative patterns. C# has key features to help with those patterns, including the ability to write unnamed functions on the fly that “capture” variables (*lambda expressions*), and the ability to perform list or reactive programming via *query expressions*. C# also makes it easy to

define read-only fields and properties for writing *immutable* (read-only) types.

## Type Safety

C# is primarily a *type-safe* language, meaning that instances of types can interact only through protocols they define, thereby ensuring each type's internal consistency. For instance, C# prevents you from interacting with a *string* type as though it were an *integer* type.

More specifically, C# supports *static typing*, meaning that the language enforces type safety at *compile time*. This is in addition to type safety being enforced at *runtime*.

Static typing eliminates a large class of errors before a program is even run. It shifts the burden away from runtime unit tests onto the compiler to verify that all the types in a program fit together correctly. This makes large programs much easier to manage, more predictable, and more robust. Furthermore, static typing allows tools such as IntelliSense in Visual Studio to help you write a program, because it knows for a given variable what type it is, and hence what methods you can call on that variable. Such tools can also identify everywhere in your program that a variable, type, or method is used, allowing for reliable refactoring.

### NOTE

C# also allows parts of your code to be dynamically typed via the `dynamic` keyword. However, C# remains a predominantly statically typed language.

C# is also called a *strongly typed language* because its type rules are strictly enforced (whether statically or at runtime). For instance, you cannot call a function that's designed to accept an integer with a floating-point number, unless you first *explicitly* convert the floating-point number to an integer. This helps prevent mistakes.

## Memory Management

C# relies on the runtime to perform automatic memory management. The Common Language Runtime has a garbage collector that executes as part of your program, reclaiming memory for objects that are no longer referenced. This frees programmers from explicitly deallocating the memory for an object, eliminating the problem of incorrect pointers encountered in languages such as C++.

C# does not eliminate pointers: it merely makes them unnecessary for most programming tasks. For performance-critical hotspots and interoperability, pointers and explicit memory allocation are permitted in blocks that are marked `unsafe`.

## Platform Support

Historically, C# was used almost entirely for writing code to run on Windows platforms. However, Microsoft and other companies have since invested in other platforms:

- The *.NET Core* Framework enables web application development in Linux and macOS (as well as Windows).

- *Xamarin* enables mobile app development for iOS and Android.
- *Blazor* compiles C# to web assembly that can run in a browser.

And on the Windows platform:

- *.NET Core 3* enables rich-client and web application development on Windows 7 to 10.
- *Universal Windows Platform* (UWP) supports Windows 10 desktop and devices such as Xbox, Surface Hub, and HoloLens.

## C# and the Common Language Runtime

C# depends on a *Common Language Runtime* (CLR), which provides essential runtime services such as automatic memory management and exception handling. (The word *common* refers to the fact that the same runtime can be shared by other *managed* programming languages, such as F#, Visual Basic, and Managed C++.)

C# is called a *managed language* because it compiles source code into managed code, which is represented in *Intermediate Language* (IL). The CLR converts the IL into the native code of the machine, such as X86 or X64, usually just prior to execution. This is referred to as Just-In-Time (JIT) compilation. Ahead-of-time compilation is also available to improve startup time with large assemblies or resource-constrained devices (and to satisfy iOS app store rules when developing with Xamarin).

The container for managed code is called an *assembly*. An assembly contains not only IL, but type information (*metadata*). The presence of metadata allows assemblies to reference types in other assemblies without needing additional files.

### NOTE

You can examine and disassemble the contents of an assembly with Microsoft's *ildasm* tool. And with tools such as ILSpy or JetBrains dotPeek, you can go further and decompile the IL to C#. Because IL is higher-level than native machine code, the decompiler can do quite a good job of reconstructing the original C#.

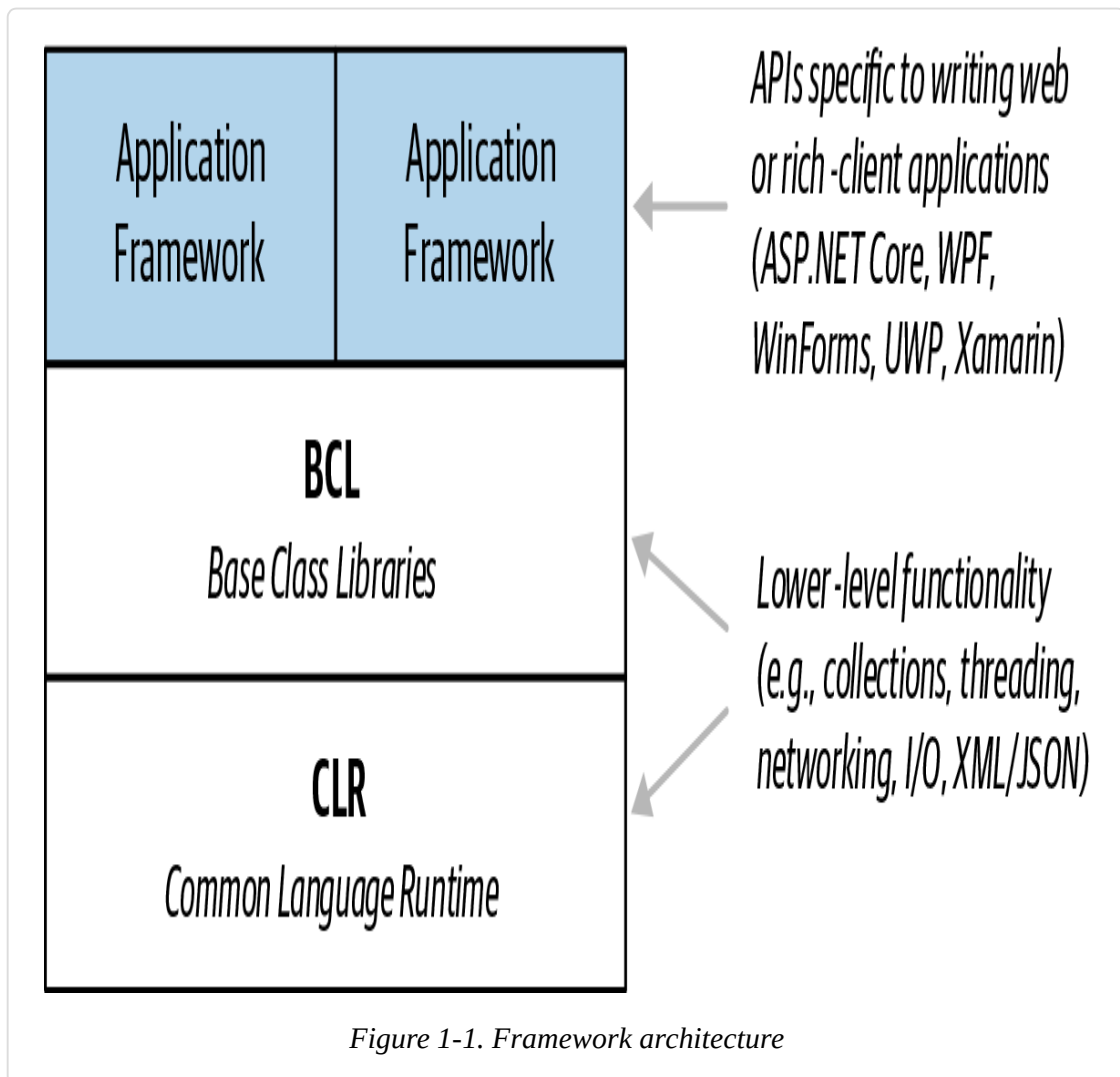
A program can query its own metadata (*reflection*) and even generate new IL at runtime (*Reflection.Emit*).

## Frameworks and Base Class Libraries

A CLR does not ship on its own, but as part of a *framework* that includes a standard set of assemblies. When writing an application, you *target* a particular framework, which means that your application uses and depends on the functionality that the framework provides. Your choice of framework also determines which platforms your application will support.

A framework comprises three layers, as illustrated in [Figure 1-1](#). The *Base Class Libraries* (BCL) sit atop the CLR, providing features useful to any kind of application (such as collections, XML/JSON, input/output [I/O], networking, serialization, and parallel

programming). Sitting atop the BCL are *application framework* layers, which provide the APIs for a user interface paradigm (such as ASP.NET Core for a web application, or Windows Presentation Foundation [WPF] for a rich-client application). A command-line program does not require an application layer.



When C# was first released in 2000, there was just the Microsoft .NET Framework. Now there are four major framework choices:

.NET Core

Modern open source framework for writing web and console applications that run on Windows, Linux, and macOS—and rich-client applications that run on Windows 7 through 10 (with .NET Core 3+). You can install multiple versions of .NET Core side by side, and applications can be *self-contained*, so as not to require a .NET Core installation.

## UWP

For writing immersive touch-first applications that run on Windows 10 desktop and devices (Xbox, Surface Hub, and HoloLens). UWP apps are sandboxed and ship via the Windows Store. UWP is preinstalled with Windows 10.

## Mono + Xamarin

Open source framework for writing mobile apps that run on iOS and Android.

## .NET Framework (superseded by .NET Core 3)

For writing web and rich-client applications that target Windows desktop/server. No major new releases are planned, although Microsoft will continue to support and maintain the current 4.8 release due to the wealth of existing applications. .NET Framework is preinstalled in Windows and supports C# 7.3 and earlier.

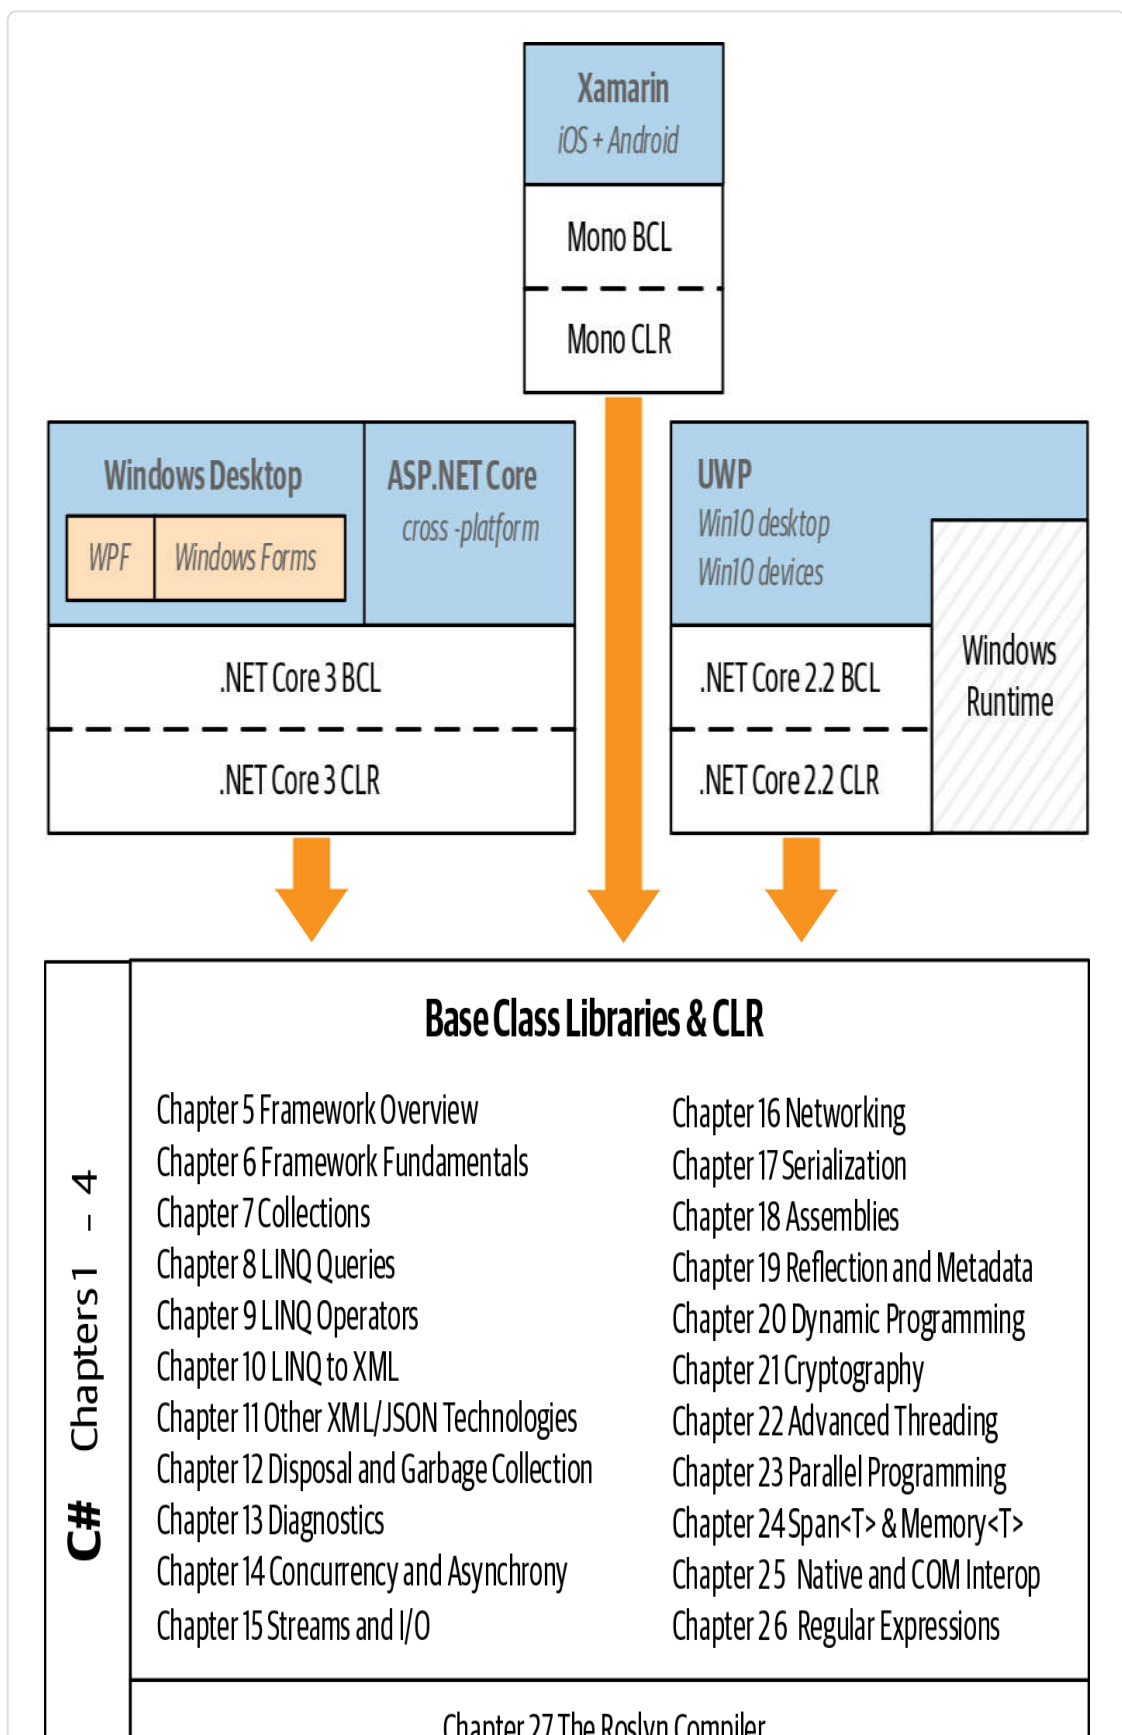
Although each of these frameworks differ in their platform support and intended uses, they all expose a similar CLR and BCL.

### NOTE

You can take advantage of this commonality and write class libraries that work across multiple frameworks—see [“.NET Standard”](#) in [Chapter 5](#).



This book focuses on C# and the core functionality of the CLR and BCL, as shown in [Figure 1-2](#). Even though the main emphasis is on .NET Core 3, we also cover some of the Windows Runtime types for UWP applications that provide functionality in parallel to the BCL.



*Figure 1-2. Topics covered in this book—the application frameworks (shown in gray) are not covered*

## Legacy and Niche Frameworks

The following frameworks are still available to support older platforms:

- Windows Runtime for Windows 8/8.1 (now UWP)
- Microsoft XNA for game development (now UWP)
- .NET Core 1.x and 2.x (for web and command-line applications only)

There are also the following niche frameworks:

- The .NET Micro Framework is for running .NET code on highly resource-constrained embedded devices (under one megabyte).
- Mono (upon which Xamarin sits) also has an application layer to develop cross-platform desktop “Windows Forms” applications on Linux, macOS, and Windows. Not all features are supported or work fully. (Another option for cross-platform user interface [UI] development is *Avalonia*, which is a WPF-inspired library that runs atop .NET Core and .NET Framework.)
- Unity is a game development platform that allows game logic to be scripted with C#.

It's also possible to run managed code within SQL Server. With SQL Server CLR integration, you can write custom functions, stored procedures, and aggregations in C# and then call them from SQL. This works in conjunction with .NET Framework and a special “hosted” CLR that enforces a sandbox to protect the integrity of the SQL Server process.

## Windows Runtime

C# also interoperates with *Windows Runtime* (WinRT) technology. WinRT means two things:

- A language-neutral object-oriented execution interface supported in Windows 8 and above
- A set of libraries baked into Windows 8 and above that support this execution interface

### NOTE

Somewhat confusingly, the term *WinRT* has historically been used to mean two more things:

- The predecessor to UWP; that is, the development platform for writing Store apps for Windows 8/8.1, sometimes called “Metro” or “Modern”
- The defunct mobile operating system for RISC-based tablets (“Windows RT”) that Microsoft released in 2011

By *execution interface*, we mean a protocol for calling code that's (potentially) written in another language. Microsoft Windows has

historically provided a primitive execution interface in the form of low-level C-style function calls comprising the Win32 API.

WinRT is much richer. In part, it is an enhanced version of Component Object Model (COM) that supports .NET, C++, and JavaScript. Unlike Win32, it's object oriented and has a relatively rich type system. This means that referencing a WinRT library from C# feels much like referencing a .NET library—you might not even be aware that you're using WinRT.

The WinRT libraries in Windows 10 form an essential part of the UWP platform (UWP relies on both WinRT and .NET Core libraries). If you're targeting the standard .NET Core platform, referencing the Windows 10 WinRT libraries is optional and can be useful if you need to access Windows 10–specific features not otherwise covered in .NET Core.

The WinRT libraries in Windows 10 support the UWP UI for writing immersive touch-first applications. They also support mobile device–specific features such as sensors, text messaging, and so on (the new functionality of Windows 8, 8.1, and 10 has been exposed through WinRT rather than Win32). WinRT libraries also provide file I/O tailored to work well within the UWP sandbox.

What distinguishes WinRT from ordinary COM is that WinRT *projects* its libraries into a multitude of languages, namely C#, Visual Basic, C++, and JavaScript, so that each language sees WinRT types (almost) as though they were written especially for it. For example, WinRT will adapt capitalization rules to suit the standards of the

target language and will even remap some functions and interfaces. WinRT assemblies also ship with rich *metadata* in *.winmd* files, which have the same format as .NET assembly files, allowing transparent consumption without special ritual; this is why you might be unaware that you're using WinRT rather than .NET types, aside from namespace differences. Another clue is that WinRT types are subject to COM-style restrictions; for instance, they offer limited support for inheritance and generics.

In C#, you not only can consume WinRT libraries, you can also write your own (and call them from a JavaScript application).

## A Brief History of C#

The following is a reverse chronology of the new features in each C# version, for the benefit of readers who are already familiar with an older version of the language.

### What's New in C# 8.0

C# 8.0 ships with *Visual Studio 2019*.

#### INDICES AND RANGES

*Indices and ranges* simplify working with elements or portions of an array (or the low-level types `Span<T>` and `ReadOnlySpan<T>`).

Indices let you refer to elements relative to the *end* of an array by using the `^` operator. `^1` refers to the last element, `^2` refers to the second-to-last element, and so on:

```
char[] vowels = new char[] { 'a', 'e', 'i', 'o', 'u' };
char lastElement = vowels [ ^1 ];    // 'u'
char secondToLast = vowels [ ^2 ];   // 'o'
```

Ranges let you “slice” an array by using the `..` operator:

```
char[] firstTwo = vowels [ ..2 ];    // 'a', 'e'
char[] lastThree = vowels [ 2.. ];   // 'i', 'o', 'u'
char[] middleOne = vowels [ 2..3 ];  // 'i'
char[] lastTwo = vowels [ ^2.. ];    // 'o', 'u'
```

C# implements indexes and ranges with the help of the `Index` and `Range` types:

```
Index last = ^1;
Range firstTwoRange = 0..2;
char[] firstTwo = vowels [ firstTwoRange ];    // 'a', 'e'
```

You can support indices and ranges in your own classes by defining an indexer with a parameter type of `Index` or `Range`:

```
class Sentence
{
    string[] words = "The quick brown fox".Split();

    public string this [Index index] => words [index];
    public string[] this [Range range] => words [range];
}
```

For more information, see [“Indices and Ranges \(C# 8\)”](#) in Chapter 2.

## NULL-COALESCING ASSIGNMENT

The `??=` operator assigns a variable only if it's null. Instead of this:

```
if (s == null) s = "Hello, world";
```

you can now write this:

```
s ??= "Hello, world";
```

## USING DECLARATIONS

If you omit the brackets and statement block following a `using` statement, it becomes a *using declaration*. The resource is then disposed when execution falls outside the *enclosing* statement block:

```
if (File.Exists ("file.txt"))
{
    using var reader = File.OpenText ("file.txt");
    Console.WriteLine (reader.ReadLine());
    ...
}
```

In this case, `reader` will be disposed when execution falls outside the `if` statement block.

## READONLY MEMBERS

C# 8 lets you apply the `readonly` modifier to a struct's *functions*, ensuring that if the function attempts to modify any field, a compile-time error is generated:



```
struct Point
{
    public int X, Y;
    public readonly void ResetX() => X = 0;    // Error!
}
```

If a **readonly** function calls a non-**readonly** function, the compiler generates a warning (and defensively copies the struct to avoid the possibility of a mutation).

## STATIC LOCAL METHODS

Adding the **static** modifier to a local method prevents it from seeing the local variables and parameters of the enclosing method. This helps to reduce coupling as well as enabling the local method to declare variables as it pleases, without risk of colliding with those in the containing method.

## DEFAULT INTERFACE MEMBERS

C# 8 lets you add a default implementation to an interface member, making it optional to implement:

```
interface ILogger
{
    void Log (string text) => Console.WriteLine (text);
}
```

This means that you can add a member to an interface without breaking implementations. Default implementations must be called explicitly through the interface:

```
((ILogger)new Logger()).Log ("message");
```

Interfaces can also define static members (including fields), which can be accessed from code inside default implementations:

```
interface ILogger
{
    void Log (string text) => Console.WriteLine (Prefix + text);
    static string Prefix = "";
}
```

or from outside the interface:

```
ILogger.Prefix = "File log: ";
```

unless restricted via an accessibility modifier on the static interface member (such as `private`, `protected`, or `internal`). Instance fields are prohibited.

For more details, see [“Default Interface Members \(C# 8\)”](#) in [Chapter 3](#).

## SWITCH EXPRESSIONS

From C# 8, you can use `switch` in the context of an *expression*:

```
string cardName = cardNumber switch    // assuming cardNumber is an
int
{
    13 => "King",
```

```

12 => "Queen",
11 => "Jack",
_ => "Pip card"    // equivalent to 'default'
};

```

For more examples, see “[switch expressions \(C# 8\)](#)” in [Chapter 2](#).

## TUPLE, POSITIONAL, AND PROPERTY PATTERNS

C# 8 supports three new patterns, mostly for the benefit of `switch` statements/expressions (see “[Patterns](#)” in [Chapter 4](#)). *Tuple patterns* let you switch on multiple values:

```

int cardNumber = 12; string suite = "spades";
string cardName = (cardNumber, suite) switch
{
    (13, "spades") => "King of spades",
    (13, "clubs") => "King of clubs",
    ...
};

```

*Positional patterns* allow a similar syntax for objects that expose a deconstructor, and *property patterns* let you match on an object’s properties. You can use all of the patterns both in switches and by the `is` operator. The following example uses a property pattern to test whether `obj` is a string with a length of 4:

```

if (obj is string { Length:4 }) ...

```

## NULLABLE REFERENCE TYPES

Whereas *nullable value types* bring nullability to value types, *nullable reference types* do the opposite and bring (a degree of) *non-nullability* to reference types, with the purpose of helping to avoid `NullReferenceExceptions`. Nullable reference types introduce a level of safety that's enforced purely by the compiler in the form of warnings or errors when it detects code that's at risk of generating a `NullReferenceException`.

Nullable reference types can be enabled either at the project level (via the `Nullable` element in the `.csproj` project file) or in code (via the `#nullable` directive). After it's enabled, the compiler makes non-nullability the default: if you want a reference type to accept nulls, you must apply the `?` suffix to indicate a *nullable reference type*:

```
#nullable enable    // Enable nullable reference types from this point
on

string s1 = null;   // Generates a compiler warning! (s1 is non-
nullable)
string? s2 = null; // OK: s2 is nullable reference type
```

Uninitialized fields also generate a warning (if the type is not marked as nullable), as does dereferencing a nullable reference type, if the compiler thinks a `NullReferenceException` might occur:

```
void Foo (string? s) => Console.Write (s.Length); // Warning
(.Length)
```

To remove the warning, you can use the *null-forgiving operator* (`!`):

```
void Foo (string? s) => Console.Write (s!.Length);
```

For a full discussion, see [“Nullable Reference Types \(C# 8\)”](#) in [Chapter 4](#).

## ASYNCHRONOUS STREAMS

Prior to C# 8, you could use `yield return` to write an *iterator*, or `await` to write an *asynchronous function*. But you couldn't do both and write an iterator that awaits, yielding elements asynchronously. C# 8 fixes this through the introduction of *asynchronous streams*:

```
async IEnumerable<int> RangeAsync (  
    int start, int count, int delay)  
{  
    for (int i = start; i < start + count; i++)  
    {  
        await Task.Delay (delay);  
        yield return i;  
    }  
}
```

The `await foreach` statement consumes an asynchronous stream:

```
await foreach (var number in RangeAsync (0, 10, 100))  
    Console.WriteLine (number);
```

For more information, see [“Asynchronous Streams \(C# 8\)”](#) in [Chapter 14](#).

## What's New in C# 7.x

C# 7 shipped with Visual Studio 2017.

### C# 7.3

C# 7.3 made minor improvements to existing features, such as enabling use of the equality operators with tuples, improved overload resolution, and the ability to apply attributes to the backing fields of automatic properties:

```
[field:NonSerialized]
public int MyProperty { get; set; }
```

C# 7.3 also built on C# 7.2's advanced low-allocation programming features, with the ability to reassign *ref locals*, no requirement to pin when indexing *fixed* fields, and field initializer support with *stackalloc*:

```
int* pointer = stackalloc int[] {1, 2, 3};
Span<int> arr = stackalloc [] {1, 2, 3};
```

Notice that stack-allocated memory can be assigned directly to a `Span<T>`. We describe spans—and why you would use them—in [Chapter 24](#).

### C# 7.2

C# 7.2 added a new `private protected` modifier (the *intersection* of `internal` and `protected`), the ability to follow named arguments with positional ones when calling methods, and `readonly` structs. A `readonly` struct enforces that all fields are `readonly`, to aid in

declaring intent and to allow the compiler more optimization freedom:

```
readonly struct Point
{
    public readonly int X, Y;    // X and Y must be readonly
}
```

C# 7.2 also added specialized features to help with micro-optimization and low-allocation programming: see [“The in modifier”](#), [“Ref Locals”](#), and [“Ref Returns”](#) in Chapter 2, and [“Ref Structs”](#) in Chapter 3.

## C# 7.1

From C# 7.1, you can omit the type when using the `default` keyword, if the type can be inferred:

```
decimal number = default;    // number is decimal
```

C# 7.1 also relaxed the rules for `switch` statements (so that you can pattern-match on generic type parameters), allowed a program’s `Main` method to be asynchronous, and allowed tuple element names to be inferred:

```
var now = DateTime.Now;
var tuple = (now.Hour, now.Minute, now.Second);
```

## NUMERIC LITERAL IMPROVEMENTS

Numeric literals in C# 7 can include underscores to improve readability. These are called *digit separators* and are ignored by the compiler:

```
int million = 1_000_000;
```

*Binary literals* can be specified with the `0b` prefix:

```
var b = 0b1010_1011_1100_1101_1110_1111;
```

## OUT VARIABLES AND DISCARDS

C# 7 makes it easier to call methods that contain `out` parameters. First, you can now declare *out variables* on the fly (see [“Out variables and discards”](#) in Chapter 2):

```
bool successful = int.TryParse ("123", out int result);  
Console.WriteLine (result);
```

And when calling a method with multiple `out` parameters, you can *discard* ones you’re uninterested in with the underscore character:

```
SomeBigMethod (out _, out _, out _, out int x, out _, out _, out  
_);  
Console.WriteLine (x);
```

## TYPE PATTERNS AND PATTERN VARIABLES

You can also introduce variables on the fly with the `is` operator. These are called *pattern variables* (see [“Introducing a pattern](#)



variable” in Chapter 3):

```
void Foo (object x)
{
    if (x is string s)
        Console.WriteLine (s.Length);
}
```

The `switch` statement also supports type patterns, so you can switch on *type* as well as constants (see “Switching on types” in Chapter 2). You can specify conditions with a `when` clause and also switch on the `null` value:

```
switch (x)
{
    case int i:
        Console.WriteLine ("It's an int!");
        break;
    case string s:
        Console.WriteLine (s.Length);    // We can use the s variable
        break;
    case bool b when b == true:           // Matches only when b is true
        Console.WriteLine ("True");
        break;
    case null:
        Console.WriteLine ("Nothing");
        break;
}
```

## LOCAL METHODS

A *local method* is a method declared within another function (see “Local methods” in Chapter 3):

```

void WriteCubes()
{
    Console.WriteLine (Cube (3));
    Console.WriteLine (Cube (4));
    Console.WriteLine (Cube (5));

    int Cube (int value) => value * value * value;
}

```

Local methods are visible only to the containing function and can capture local variables in the same way that lambda expressions do.

## MORE EXPRESSION-BODIED MEMBERS

C# 6 introduced the expression-bodied *fat-arrow* syntax for methods, read-only properties, operators, and indexers. C# 7 extends this to constructors, read/write properties, and finalizers:

```

public class Person
{
    string name;

    public Person (string name) => Name = name;

    public string Name
    {
        get => name;
        set => name = value ?? "";
    }

    ~Person () => Console.WriteLine ("finalize");
}

```

## DECONSTRUCTORS

C# 7 introduces the *destructor* pattern (see “[Deconstructors](#)” in [Chapter 3](#)). Whereas a constructor typically takes a set of values (as parameters) and assigns them to fields, a *destructor* does the reverse and assigns fields back to a set of variables. We could write a destructor for the `Person` class in the preceding example as follows (exception-handling aside):

```
public void Deconstruct (out string firstName, out string lastName)
{
    int spacePos = name.IndexOf ( ' ');
    firstName = name.Substring (0, spacePos);
    lastName = name.Substring (spacePos + 1);
}
```

Deconstructors are called with the following special syntax:

```
var joe = new Person ("Joe Bloggs");
var (first, last) = joe;           // Deconstruction
Console.WriteLine (first);        // Joe
Console.WriteLine (last);         // Bloggs
```

## TUPLES

Perhaps the most notable improvement to C# 7 is explicit *tuple* support (see “[Tuples](#)” in [Chapter 4](#)). Tuples provide a simple way to store a set of related values:

```
var bob = ("Bob", 23);
Console.WriteLine (bob.Item1);    // Bob
Console.WriteLine (bob.Item2);    // 23
```

C#'s new tuples are syntactic sugar for using the `System.ValueTuple<...>` generic structs. But thanks to compiler magic, tuple elements can be named:

```
var tuple = (name:"Bob", age:23);
Console.WriteLine (tuple.name);    // Bob
Console.WriteLine (tuple.age);     // 23
```

With tuples, functions can return multiple values without resorting to out parameters or extra type baggage:

```
static (int row, int column) GetFilePosition() => (3, 10);

static void Main()
{
    var pos = GetFilePosition();
    Console.WriteLine (pos.row);    // 3
    Console.WriteLine (pos.column); // 10
}
```

Tuples implicitly support the deconstruction pattern, so you can easily *deconstruct* them into individual variables:

```
static void Main()
{
    (int row, int column) = GetFilePosition(); // Creates 2 local
    variables
    Console.WriteLine (row);    // 3
    Console.WriteLine (column); // 10
}
```

## THROW EXPRESSIONS

Prior to C# 7, `throw` was always a statement. Now it can also appear as an expression in expression-bodied functions:

```
public string Foo() => throw new NotImplementedException();
```

A `throw` expression can also appear in a ternary conditional expression:

```
string Capitalize (string value) =>
    value == null ? throw new ArgumentException ("value") :
    value == "" ? "" :
    char.ToUpper (value[0]) + value.Substring (1);
```

## What's New in C# 6.0

C# 6.0, which shipped with Visual Studio 2015, features a new-generation compiler, completely written in C#. Known as project “Roslyn,” the new compiler exposes the entire compilation pipeline via libraries, allowing you to perform code analysis on arbitrary source code (see [Chapter 27](#)). The compiler itself is open source, and the source code is available [on GitHub](#).

In addition, C# 6.0 features several minor but significant enhancements, aimed primarily at reducing code clutter.

The *null-conditional* (“Elvis”) operator (see [“Null Operators”](#) in [Chapter 2](#)) avoids having to explicitly check for null before calling a method or accessing a type member. In the following example, `result` evaluates to null instead of throwing a `NullReferenceException`:

```
System.Text.StringBuilder sb = null;  
string result = sb?.ToString();    // result is null
```

*Expression-bodied functions* (see “[Methods](#)” in [Chapter 3](#)) allow methods, properties, operators, and indexers that comprise a single expression to be written more tersely, in the style of a lambda expression:

```
public int TimesTwo (int x) => x * 2;  
public string SomeProperty => "Property value";
```

*Property initializers* ([Chapter 3](#)) let you assign an initial value to an automatic property:

```
public DateTime TimeCreated { get; set; } = DateTime.Now;
```

Initialized properties can also be read-only:

```
public DateTime TimeCreated { get; } = DateTime.Now;
```

Read-only properties can also be set in the constructor, making it easier to create immutable (read-only) types.

*Index initializers* ([Chapter 4](#)) allow single-step initialization of any type that exposes an indexer:

```
var dict = new Dictionary<int,string>()  
{  
    [3] = "three",
```

```
[10] = "ten"
};
```

*String interpolation* (see “[String Type](#)” in [Chapter 2](#)) offers a succinct alternative to `string.Format`:

```
string s = $"It is {DateTime.Now.DayOfWeek} today";
```

*Exception filters* (see “[try Statements and Exceptions](#)” in [Chapter 4](#)) let you apply a condition to a `catch` block:

```
string html;
try
{
    html = new WebClient().DownloadString ("http://asef");
}
catch (WebException ex) when (ex.Status ==
    WebExceptionStatus.Timeout)
{
    ...
}
```

The `using static` (see “[Namespaces](#)” in [Chapter 2](#)) directive lets you import all the static members of a type so that you can use those members unqualified:

```
using static System.Console;
...
WriteLine ("Hello, world"); // WriteLine instead of Console.WriteLine
```

The `nameof` ([Chapter 3](#)) operator returns the name of a variable, type, or other symbol as a string. This avoids breaking code when you rename a symbol in Visual Studio:

```
int capacity = 123;
string x = nameof (capacity);    // x is "capacity"
string y = nameof (Uri.Host);    // y is "Host"
```

And finally, you're now allowed to `await` inside `catch` and `finally` blocks.

## What's New in C# 5.0

C# 5.0's big new feature was support for *asynchronous functions* via two new keywords, `async` and `await`. Asynchronous functions enable *asynchronous continuations*, which make it easier to write responsive and thread-safe rich-client applications. They also make it easy to write highly concurrent and efficient I/O-bound applications that don't tie up a thread resource per operation.

We cover asynchronous functions in detail in [Chapter 14](#).

## What's New in C# 4.0

C# 4.0 introduced four major enhancements:

- *Dynamic binding* ([Chapters 4](#) and [20](#)) defers *binding*—the process of resolving types and members—from compile time to runtime and is useful in scenarios that would otherwise require complicated reflection code. Dynamic binding is also useful when interoperating with dynamic languages and COM components.



- *Optional parameters* (Chapter 2) allow functions to specify default parameter values so that callers can omit arguments, and *named arguments* allow a function caller to identify an argument by name rather than position.
- *Type variance* rules were relaxed in C# 4.0 (Chapters 3 and 4), such that type parameters in generic interfaces and generic delegates can be marked as *covariant* or *contravariant*, allowing more natural type conversions.
- *COM interoperability* (Chapter 25) was enhanced in C# 4.0 in three ways. First, arguments can be passed by reference without the `ref` keyword (particularly useful in conjunction with optional parameters). Second, assemblies that contain COM interop types can be *linked* rather than *referenced*. Linked interop types support type equivalence, avoiding the need for *Primary Interop Assemblies* and putting an end to versioning and deployment headaches. Third, functions that return COM-Variant types from linked interop types are mapped to `dynamic` rather than `object`, eliminating the need for casting.

## What's New in C# 3.0

The features added to C# 3.0 were mostly centered on *Language-Integrated Query* (LINQ) capabilities. LINQ enables queries to be written directly within a C# program and checked *statically* for correctness, and query both local collections (such as lists or XML documents) or remote data sources (such as a database). The C# 3.0 features added to support LINQ comprised implicitly typed local variables, anonymous types, object initializers, lambda expressions, extension methods, query expressions, and expression trees.

*Implicitly typed local variables* (`var` keyword, [Chapter 2](#)) let you omit the variable type in a declaration statement, allowing the compiler to infer it. This reduces clutter as well as allowing *anonymous types* ([Chapter 4](#)), which are simple classes created on the fly that are commonly used in the final output of LINQ queries. You can also implicitly type arrays ([Chapter 2](#)).

*Object initializers* ([Chapter 3](#)) simplify object construction by allowing you to set properties inline after the constructor call. Object initializers work with both named and anonymous types.

*Lambda expressions* ([Chapter 4](#)) are miniature functions created by the compiler on the fly; they are particularly useful in “fluent” LINQ queries ([Chapter 8](#)).

*Extension methods* ([Chapter 4](#)) extend an existing type with new methods (without altering the type’s definition), making static methods feel like instance methods. LINQ’s query operators are implemented as extension methods.

*Query expressions* ([Chapter 8](#)) provide a higher-level syntax for writing LINQ queries that can be substantially simpler when working with multiple sequences or range variables.

*Expression trees* ([Chapter 8](#)) are miniature code Document Object Models (DOMs) that describe lambda expressions assigned to the special type `Expression<TDelegate>`. Expression trees make it possible for LINQ queries to execute remotely (e.g., on a database

server) because they can be introspected and translated at runtime (e.g., into a SQL statement).

C# 3.0 also added automatic properties and partial methods.

*Automatic properties* (Chapter 3) cut the work in writing properties that simply `get/set` a private backing field by having the compiler do that work automatically. *Partial methods* (Chapter 3) let an autogenerated partial class provide customizable hooks for manual authoring which “melt away” if unused.

## **What’s New in C# 2.0**

The big new features in C# 2 were generics (Chapter 3), nullable value types (Chapter 4), iterators (Chapter 4), and anonymous methods (the predecessor to lambda expressions). These features paved the way for the introduction of LINQ in C# 3.

C# 2 also added support for partial classes, static classes, and a host of minor and miscellaneous features such as the namespace alias qualifier, friend assemblies, and fixed-size buffers.

The introduction of generics required a new CLR (CLR 2.0), because generics maintain full type fidelity at runtime.

# Chapter 2. C# Language Basics

---

In this chapter, we introduce the basics of the C# language.

## NOTE

All programs and code snippets in this and the following two chapters are available as interactive samples in LINQPad. Working through these samples in conjunction with the book accelerates learning in that you can edit the samples and instantly see the results without needing to set up projects and solutions in Visual Studio.

To download them in *LINQPad*, click the Samples tab, and then click “Download more samples.”

## A First C# Program

Following is a program that multiplies 12 by 30 and prints the result, 360, to the screen. The double forward slash indicates that the remainder of a line is a *comment*:

```
using System;                                // Importing namespace

class Test                                   // Class declaration
{
    static void Main()                       // Method declaration
    {
```